

Instruction Set Architecture

Isaac Pérez Andrade

ITESM Guadalajara
Department of Engineering & Architecture
Department of Computer Science

August - December 2020



Instruction Set Architecture

Instruction Set Architecture (ISA)

- It is an **abstract** concept which defines the portion of a computer that is visible to both the programmer and the compiler.
- It is part of the link between an application (something a human does such as video recording, playing music, editing a spreadsheet, etc) and the physical layer of the computer, *i.e.* Hardware (HW).

Instruction Set Architecture

ISA vs μA

- ISA **theoretically** describes how a computer executes its programs.
- It describes:
 - The fundamental operations, which are simply referred to as *instructions*, that the computer can execute.
 - How these instructions are executed.
 - The semantics and rules required for the interaction of the different building blocks of a computer.

Instruction Set Architecture

ISA vs μA

- Overall, ISA provides valuable information to the programmer.
 - Is a computer stack-, accumulator- or register-based?
 - Does the computer have memory? Does it have registers?
 - How many steps, *i.e.* clock cycles, does it take to execute instructions?
 - Where are operands fetched from?
 - Where is the result stored?
 - How big are data types?

Microarchitecture (μA)

- μA is more closely related to the **physical** implementation of a design, *i.e.*, μA determines how the ISA is implemented in HW.
 - For example, it describes which building are necessary in order to model a Microprocessor (μP) and how these building blocks are connected with each other.

Instruction Set Architecture

ISA vs μA

- The same ISA may be physically implemented in a variety of μA s.
 - For example, one ISA could be implemented by different HW approaches and vendors such as Intel, ARM or AMD, and all three could have different performances.
- A naive adder example:
 - ISA specifies data width as 64-bits.
 - μA defines the adder as ripple-carry, carry-lookahead, carry-save, carry-select, etc.

Instruction Set Architecture

ISA vs μA

- The goal of a processor designer is to evaluate the different trade-offs between ISA and μA in order to find a Pareto optimal system.
 - Power consumption.
 - Latency - How long does it take to complete a task.
 - Throughput - How many tasks can be completed in a given time.
 - Chip area.

Instruction Set Architecture

Same ISA, different μA - 45 nm technology

- x86 ISA.
- Quad Core.
- 2.6 GHz.
- 125 W.




Figure 1: AMD Phenom X4

- x86 ISA.
- Dual Core.
- 1.6 GHz.
- 2 W.




Figure 2: Intel Atom

Instruction Set Architecture

Different ISA, different μ A- 45 nm technology

- x86 ISA.
 - Quad Core.
 - 2.6 GHz.
 - 125 W.
- Power ISA.
 - Octa Core.
 - 4.25 GHz.
 - 200 W.




Figure 3: AMD Phenom X4




Figure 4: IBM Power7

ISA characteristics

ISA characteristics

- Type and size of instructions.
- Type and size of operands.
- Instruction encoding.
- Addressing modes.
- Registers

Registers

ISA characteristics

Registers

This is list of registers commonly found in $\mu\mathbf{P}$. Note that generally, $\mu\mathbf{P}$ s do not implement every single register in this list.

- **Program Counter (PC)**. Also called **IP**, points to the memory address of the next instruction to be executed.
- **Register File (RF)**. Set of registers used to store data.
- **Stack Pointer (SP)**. Points to the next location in the stack. Used in PUSH and POP operations.

Basic registers of a computer

- **Instruction Register (IR)**. Holds the instruction currently being executed.
- **Memory Address Register (MAR)**. Also called **Address Register (AR)**, points to the memory address to/from which data is stored/fetched to/from.
- **Instruction Memory (IM)**. Memory that stores the instructions that the μ P will execute.

Basic registers of a computer

- **Accumulator (ACC)**. Holds the result of arithmetic and logic operations.
- **Data Memory (DM)**. Also called **Data Register (DR)**, holds operand(s) to be used in arithmetic and logic operations.
- **General Purpose Register (GPR)**. Registers to temporary store data or addresses.
- **Status Register (SR)**. Also called **Flag Register (FR)**, holds the special conditions of the result of arithmetic and logic operations, as well as branch and jump status. For example, indicates if a comparison resulted in an equality, if the result of an operation is zero, overflow, etc.

Basic registers of a computer

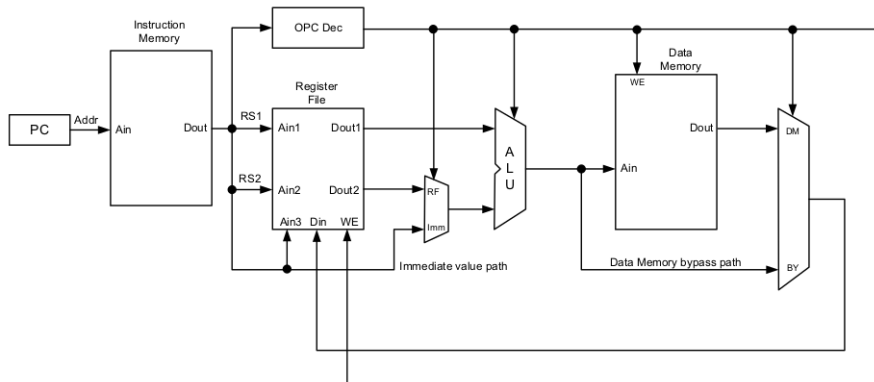


Figure 5: Basic structure of a μP .

Instructions

ISA characteristics

Instructions


- Instructions are the basic operations that a μP can understand and perform.
- The complete set of instructions that a μP may perform is called **instruction set**, **which is not the same as Instruction Set Architecture!**

ISA characteristics

Instructions

Instruction type	Example ¹
Arithmetic & logical	ADD, SUB, AND, OR
Data transfer	LOAD, SW, MOV, PUSH, POP
Conditional branch	BNE, BEQ
Unconditional jumps	JMP, JAL, CALL, RET
System	RD_INT, PRNT_CHR
Floating Point	FADD, FMULT
String	MOVSB, STR_MV, STR_CMP
Signal processing ²	ADD_ARRAY, MULT_ARRAY, FFT

¹These examples are not specific to a particular ISA.

²Typically found in Single Instruction Multiple Data (SIMD) ISAs. 

ISA characteristics

Instructions

Table 1: Intel's 80x86 top ten instructions based on five SPECint92 programs³.

Rank	Type	Distribution
1	load	22%
2	conditional branch	20%
3	compare	16%
4	store	12%
5	add	8%
6	and	6%
7	sub	5%
8	move register-register	4%
9	call	1%
10	return	1%
Total		96%

³J. L. Hennessy and D. A. Patterson, *Computer architecture: A quantitative approach*, 6th ed., p A-4, Morgan Kaufmann, 2019.

Operands and operations

ISA characteristics

Operands and operations

- Where do operands come from?
- Where are results stored?
- What is the size of the operands?
- How many steps does an instruction take?

ISA characteristics

Operands and operations

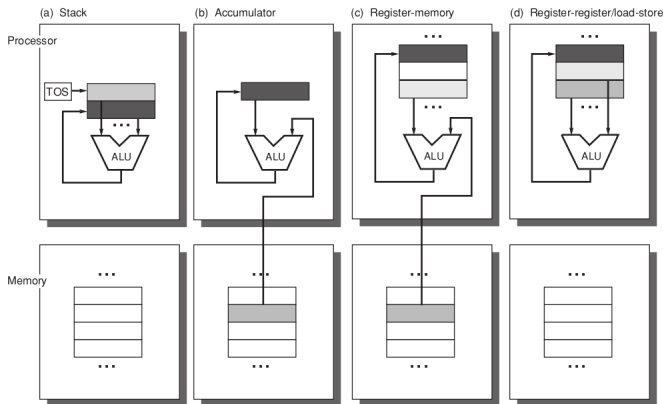



Figure 6: Operand locations for different ISAs [Figure A.1] ⁴.

⁴J. L. Hennessy and D. A. Patterson, *Computer architecture: A quantitative approach*, 6th ed., p A-4, Morgan Kaufmann, 2019.

ISA characteristics

Operands and operations



Stack-based ISA

$$C = A + B$$

Push A


Push B

Add

Pop C

ISA characteristics

Operands and operations



Accumulator-based ISA

$$C = A + B$$

Load A


Add B

Store C

ISA characteristics

Operands and operations

(c) Register-memory



Memory-Register-based ISA


$$C = A + B$$

Load	R1	A
Add	R3	R1 B
Store	R3	C

ISA characteristics

Operands and operations

(d) Register-register/load-store



Register-Register-based ISA

$$C = A + B$$

Load	R1	A	
Load	R2	B	
Add	R3	R1	R2
Store	R3	C	

Addressing Modes

ISA characteristics

Addressing modes

- How can we read/write data from/into memory?
- What types of memory exist?

ISA characteristics

Addressing modes

Table 2: Examples of addressing modes

Mode	Example	Meaning
Immediate	Add R4 , 3	$R4 \leftarrow R4 + 3$
Register	Add R4 , R3	$R4 \leftarrow R4 + R3$
Absolute (Direct)	Add R2 , (100)	$R2 \leftarrow R2 + \text{Mem}[100]$
Register indirect	Add R4 , (R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$
Indexed	Add R3 , (R1 + R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$
Displacement	Add R4 , 100(R1)	$R4 \leftarrow R4 + \text{Mem}[100+R1]$
Memory indirect	Add R1 , @(R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$

ISA characteristics

Addressing modes

Table 3: Examples of addressing modes

Mode	Example	Meaning
Autoincrement	Add R1 , (R2)+	$R1 \leftarrow R1 + \text{Mem}[R2]$ $R2 \leftarrow R2 + d$
Autodecrement	Add R1 , -(R2)	$R2 \leftarrow R2 - d$ $R1 \leftarrow R1 + \text{Mem}[R2]$
Scaled	Add R1 , 100(R2) [R3]	$R1 \leftarrow R1 + \text{Mem}[100 \cdot R2 + R3 \cdot d]$

Note that Autoincrement and Autodecrement modes, the order of + and - signs influence the order of the operations. For example, -(R2) indicates decrementing R2 before accessing to memory.

ISA characteristics

Addressing modes: Example

- Let's assume that registers R_i , $i \in [1, 4]$, and selected memory locations of a computer store the following values.

Reg	Val	Mem	Val
R1	23	7	23
R2	11	11	13
R3	7	13	31
R4	19	23	17
		34	37
		100	13
		123	29
		132	41

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Immediate	Add R4 , 3	$R4 \leftarrow R4 + 3$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Register	Add R4 , R3	$R4 \leftarrow R4 + R3$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Absolute (Direct)	Add R2 , (100)	$R2 \leftarrow R2 + \text{Mem}[100]$

Reg	Val	Mem	Val	R2 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Register indirect	Add R4 , (R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Indexed	Add R3 , (R1 + R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$

Reg	Val	Mem	Val	R3 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Displacement	Add R4 , 100(R1)	$R4 \leftarrow R4 + \text{Mem}[100+R1]$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Memory indirect	Add R1 , @(R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$

Reg	Val	Mem	Val	R1 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Autoincrement	Add R1 , (R2)+	$R1 \leftarrow R1 + \text{Mem}[R2]$ $R2 \leftarrow R2 + d$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 3$
 $R1 = ?$

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Autodecrement	Add R1 , -(R2)	$R2 \leftarrow R2 - d$ $R1 \leftarrow R1 + \text{Mem}[R2]$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 4$
 $R1 = ?$

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Scaled	Add R1 , 100(R2) [R3]	$R1 \leftarrow R1 + \text{Mem}[100+R2 + R3 *d]$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 3$
 $R1 = ?$

Instruction encoding

ISA characteristics

Instructions encoding

- In stored-program computers, instructions and data are stored in memory.
- So, how does a processor know
 - how to differentiate between operations and operands?
 - which instruction to perform?
 - which Reg or Mem locations are the operands located?
 - which Reg or Mem locations should the result be stored to?
 - how to differentiate between

$R1 \leftarrow R2 + R3$

$R1 \leftarrow R2 + \text{Mem}[R3]$

$R1 \leftarrow R2 + \text{Mem}[\text{Mem}[R3]]$

$R1 \leftarrow \text{Mem}[R2] + \text{Mem}[R3]$

- Instruction encoding is a **convention used to differentiate the various operations in a μP , as well as operations from operands.**

ISA characteristics

Instructions encoding

- Instructions are encoded using binary representation.
- Suppose we want to design a processor that can implement the following instructions.

$R1 \leftarrow R2 + R3$

$R1 \leftarrow R2 + \text{Mem}[R3]$

$R1 \leftarrow R2 + \text{Mem}[\text{Mem}[R3]]$

$R1 \leftarrow \text{Mem}[R2] + \text{Mem}[R3]$

- We could assign a binary code to each of these 4 operations.

Instruction	Binary code
$R1 \leftarrow R2 + R3$	00
$R1 \leftarrow R2 + \text{Mem}[R3]$	01
$R1 \leftarrow R2 + \text{Mem}[\text{Mem}[R3]]$	10
$R1 \leftarrow \text{Mem}[R2] + \text{Mem}[R3]$	11

- Let's try to expand this implementation.

ISA characteristics

Instructions encoding: A naive example

- Let Rd be the destination register and Rsi the source register, where $d \in [0, 3]$ and $i \in [0, 3]$.
- We could assign a binary code for each combination of Rd and Rsi in the instruction $Rd \leftarrow Rsi + Rsi$.

Instruction	Binary code	Hex code
$R0 \leftarrow R0 + R0$	00 0000	00h
$R0 \leftarrow R0 + R1$	00 0001	01h
\vdots	\vdots	\vdots
$R1 \leftarrow R2 + R3$	01 1011	1Bh
\vdots	\vdots	\vdots
$R2 \leftarrow R3 + R0$	10 1100	2Ch
\vdots	\vdots	\vdots
$R3 \leftarrow R3 + R3$	11 1111	3Fh

ISA characteristics

Instructions encoding: A naive example

- What about the μA of this encoding?




Figure 7: A naive μA for adding two registers.

ISA characteristics

Instructions encoding

- Is our previous scheme feasible?
- What's wrong with it?
- Could it be generalised?
- What about other addressing modes?
- How could we include other operations such as subtractions or jumps?
- Are all instructions represented using the same number of bits?

ISA characteristics

Instructions encoding

- We can continue to expand this scheme in order to add other operations, *e.g.*, subtraction and logical operations.
- Moreover, we can continue to include bits that represent different addressing modes.
- The ultimate goal of this, is to design an encoding feasible for all operations and addressing modes in our ISA.

ISA characteristics

Instructions encoding

- We could have a bit for selecting addition/subtraction in our previous design.

Table 4: Naive encoding for adding and subtracting two numbers.

Instruction	Binary code	Hex code
$R0 \leftarrow R0 + R0$	000 0000	00h
\vdots	\vdots	\vdots
$R1 \leftarrow R2 + R3$	001 1011	1Bh
\vdots	\vdots	\vdots
$R3 \leftarrow R3 + R3$	011 1111	3Fh
$R0 \leftarrow R0 - R0$	100 0000	10h
\vdots	\vdots	\vdots
$R1 \leftarrow R2 - R3$	101 1011	5Bh
\vdots	\vdots	\vdots
$R3 \leftarrow R3 - R3$	111 1111	7Fh

ISA characteristics


Instructions encoding

- As seen in the previous example, the source, destination and operation type may be represented with a **single** binary code.
- This concept may be further expanded for other operations and addressing modes.
- For this purpose, we may use longer binary codes, which may be divided into several fields.
- For example, we may use a field called **opcode** in order to represent the type of operation to be performed.
- Another field may represent the **source**, *i.e.*, the memory location where data to operate with is.
- Another field may represent the **destination**, *i.e.*, the memory location where the result of the operation should be stored to.


ISA characteristics

Instructions encoding

Variable encoding



Fixed encoding



Hybrid encoding


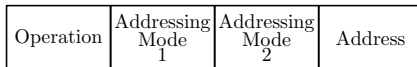


Figure 8: Generalised instruction encoding.

ISA characteristics

Instructions encoding

- **Variable encoding.**

- Supports any number of operands, with each operand having a specific addressing mode.
- The number of encoded bits varies between instructions.
- **Compact machine programs.**
- **Harder to decode.**

- **Fixed encoding.**

- Every instruction has the same number of operands, with addressing modes specified in the opcode.
- The number of encoded bits is always the same regardless of the instruction or addressing modes of the operands.
- **Easier to decode.**
- **Wasted bits in some instructions.**

- **Hybrid encoding.**

- Instructions with two different encoding lengths (16- and 32-bits, for example).

ISA characteristics

Instructions encoding

- ISAs may be classified into two main categories according to the complexity of their instruction encoding.
- **Reduced Instruction Set Computer (RISC).**
- **Complex Instruction Set Computer (CISC).**

CISC

- ISAs that perform complex operations and the instruction formats are not uniform.
- Large number of instructions available.
- Microcode approach.
 - A single instruction may be divided into several smaller instructions.
 - For example, a single instruction may perform a load from memory, an arithmetic operation and a store to memory.
- Reduced size of the compiled code due to **variable-length** encoding.
 - Shortest encodings represent the most commonly used instructions.

RISC

- ISAs that have a small number of simple, **fixed-length** instructions.
- Single-cycle instructions.
- Load-store approach.
 - Only load and store instructions are used for transferring data between registers and memory.

ISA characteristics

CISC vs RISC

```
1:  mul16:
2:      pushl   %ebp                ; 01010101
3:      movl    %esp, %ebp          ; 1000100111100101
4:      movl    8(%ebp), %ecx        ; 100010000100110100001000
5:      pushl   %ebx                ; 01010011
6:      movl    12(%ebp), %edx       ; 100010110101010100001100
7:      xorl    %ebx, %ebx           ; 0011000111011011
8:      movl    $15, %eax           ; 1011100000001111
9:      .p2align 2,,3               ; 000000000000000000000000
                                   ; 100011010111011000000000
10:  .L6:
11:      testb   $1, %dl             ; 111101101100001000000001
12:      je      .L5                 ; 0111010000000010
13:      addl    %ecx, %ebx          ; 0000000111001011
14:  .L5:
15:      sall    %ecx                ; 1101000111100001
16:      shrl    %edx                ; 1101000111101010
17:      decl    %eax                ; 01001000
18:      jns     .L6                 ; 0111100111110010
19:      movl    %ebx, %eax          ; 1000100111011000
20:      popl    %ebx                ; 01011011
21:      leave   ; 11001001
22:      ret      ; 11000011
```

Figure 9: CISC code example.

ISA characteristics

CISC vs RISC

```
1:  mul16:
2:      move    $6, $0                # 000000000000000000011000000100001
3:      li      $3, 15                # 00100100000000011000000000001111
4:  $L6:
5:      andi     $2, $5, 0x1          # 00110000101000100000000000000001
6:      addiu    $3, $3, -1           # 00100100011000111111111111111111
7:      beq      $2, $0, $L5          # 00010000010000000000000000000010
8:      srl      $5, $5, 1            # 00000000000001010010100001000010
9:      addu     $6, $6, $4           # 00000000110001000011000000100001
10: $L5:
11:      bgez     $3, $L6              # 00000100011000011111111111111010
12:      sll      $4, $4, 1            # 00000000000000100001000000100000
13:      j        $31                  # 00000011111000000000000000001000
14:      move     $2, $6              # 00000000110000000001000000100001
```

Figure 10: RISC code example.

RISC-V

ISA characteristics

RISC-V example

RISC-V main features.

- 32-bit encoding.
- 4 types of instructions.
 - R-type. Register.
 - I-type. Immediate.
 - S-type. Store, compare and branch.
 - U-type. Jump.

ISA characteristics

RISC-V example




Figure 11: RISC instruction format.

ISA characteristics

RISC encoding example

- `funct7` and `funct3` complement the opcode.
- `rd`, `rs1` and `rs2` are destination, source 1 and source 2 registers, respectively.
- `imm` is an immediate value.

ISA characteristics

RISC encoding example

Instruction format	Primary use	rd	rs1	rs2	Immediate
R-type	Register-register ALU instructions	Destination	First source	Second source	
I-type	ALU immediates Load	Destination	First source base register		Value displacement
S-type	Store Compare and branch		Base register first source	Data source to store second source	Displacement offset
U-type	Jump and link Jump and link register	Register destination for return PC	Target address for jump and link register		Target address for jump and link

Figure 12: RISC instruction description.

QUIZ

- Which of the following are affected by the instruction encoding?
 - A) The execution time of each instruction.
 - B) The μA of the processor.
 - C) Global warming.
 - D) The size of the compiled program.
 - E) All of the above.
 - F) None of the above.

QUIZ

- Which of the following are affected by the instruction encoding?
 - A) **The execution time of each instruction.**⁵
 - B) **The μ A of the processor.**
 - C) Global warming.
 - D) **The size of the compiled program.**
 - E) All of the above.
 - F) None of the above.

⁵Think about CISC and RISC differences.

Summary

Summary

- ISA is the link between applications and HW.
- μA refers to the physical implementation of the ISA.
- The same ISA can be implemented in different μA s.
- ISA encloses
 - Type and size of instructions and operands.
 - Addressing modes.
 - Instruction encoding.
- There are RISC and CISC ISAs.
- There are several trade-offs associated between ISAs and μA s, and our goal is to find a Pareto-optimal design.

Further Reading

- Read about the difference between Von-Neumann and Harvard architectures.