

Computational abstraction & Instruction Set Architecture

Isaac Pérez Andrade

ITESM Guadalajara
Department of Engineering & Architecture
Department of Computer Science

February - June 2021



The following material has been adopted and adapted from

Patterson, D. A., Hennessy, J. L., *Computer Organization and design: The hardware/software interface - ARM edition*, Morgan Kaufmann, 2017.

S. L. Harris and D. M. Harris, *Digital design and computer architecture - ARM edition*, Morgan Kaufman, 2016.

J. Yiu, *The definitive guide to ARM Cortex-M0 and Cortex-M0+ processors*, Second edition, Elsevier, 2015.

Computational abstraction

Computational abstraction

- How would you describe the operation of a computer?
- How do users communicate with Integrated Circuits (ICs) inside a computer?
 - Think about layers.
 - Between an application (Software (SW)) and the physics of a computer (Hardware (HW), ICs, transistors), there are multiple layers that communicate with each other.
 - Each layer deals with a different complexity level.
 - Can you name some of these layers?




Figure 1: Simplified view of HW and SW as hierarchical layers.

- **Applications** provide direct user interaction.
- **System** layer consists of compilers and Operative Systems (OSs).
- **HW** layer relates to the actual physics of the computer, *i.e.*, signals, voltages, shapes, *etc.*




Figure 1: Simplified view of HW and SW as hierarchical layers.

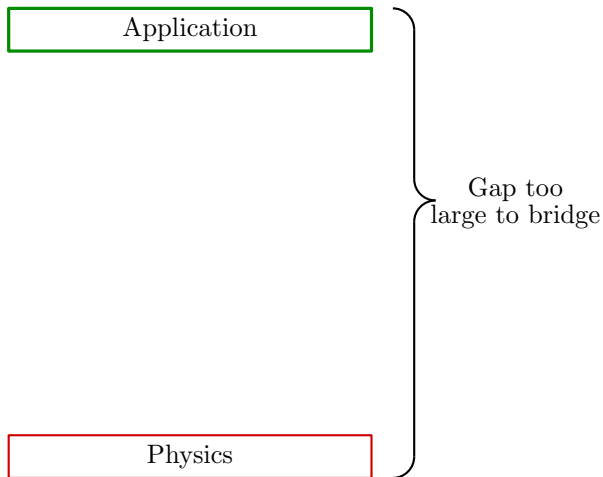
OS basic tasks.

- Provide interaction between user's programs and the HW.
- Handling basic Input/Output (IO) operations.
- Allocating storage and memory.
- Providing for protected sharing of the computer among multiple applications using it simultaneously.

Computational abstraction

- Applications such as word processors, internet browsers or media players consists of millions of lines of codes.
- Microprocessors (μ Ps) are only capable of executing extremely simple low-level instructions such as additions, comparisons, jumps, load/store from/to memory.
- Moreover, we must communicate with HW by simply using electrical signals.
- So how does a computing system perform such complex tasks using limited resources?

Computational abstraction



- **Q: How do we fill the gap between the application and the physics?**
- **A: The concept of abstraction.**
- Abstraction is no other thing than using different representations of a concept in order to deal with different levels of complexity.
- Hiding details when they are not important.
- This allows us to focus on a given complexity level and omit unnecessary details.

Computational abstraction

- Let's try to bridge the gap starting with the application layer.
- The application layer is the closest to the user.
- Application layer communicates with lower-level layers in order to instruct HW what to do.
- The following elements may be used for this purpose.
 - **High-level language.**
 - **Compiler.**
 - **Assembly language.**
 - **Assembler.**
 - **Machine language.**

Computational abstraction

- **High-level language.** Set of words and algebraic notation close to human language used to indicate a sequence of instructions.



```
#include <stdio.h>

int main () {

    int a;

    /* for loop execution */
    for( a = 1; a <= 100; a++ )
    {
        printf("%d\n",a * a);
    }

    return 0;
}
```

- **Compiler.** Software that translates high-level language to assembly language.

Computational abstraction

- **Assembly language.** Set of words, also called mnemonics or dictionaries, that symbolically represent machine instructions.

```
sub_fcd7:
    clrc                ;fcd7 81
    addcw a             ;fcd8 23
    movw ix, a          ;fcd9 e2
    mov a, @ix+0x00     ;fcda 06 00
    mov mem_009e, a     ;fcdc 45 9e
    mov a, @ix+0x01     ;fcde 06 01
    mov mem_009f, a     ;fce0 45 9f
    mov a, @ix+0x02     ;fce2 06 02
    mov mem_00a0, a     ;fce4 45 a0
    call sub_fba3        ;fce6 31 fb a3
    ret                 ;fce9 20

sub_fcea:
    movw ix, #mem_02b6  ;fcea e6 02 b6
    mov a, mem_01ef     ;fcfd 60 01 ef
    bne lab_fcf9        ;fcf0 fc 07
    mov a, #0x10        ;fcf2 04 10
    mov mem_01ef, a     ;fcf4 61 01 ef
    setb mem_0097:2     ;fcf7 aa 97

lab_fcf9:
    call sub_fd31       ;fcf9 31 fd 31
```

```
.text
main:

    # Polling keyboard input
poll_keyboard:
    lw    $t0, KEYBOARD_RX_CTRL
    beqz  $t0, poll_keyboard # 0xffff0000 = 0x00000001 ? -> Key pressed?

    # Load character from keyboard into $a0
    # $a0 contains character to be printed in syscall 11
    lbu   $a0, KEYBOARD_RX_DATA

    # Print character from keyboard and newline using syscall
    li    $v0, SYSCALL_PRINT_CHAR
    syscall

    # Polling display output
poll_display:
    lw    $t1, DISPLAY_TX_CTRL
    beqz  $t1, poll_display
    sw    $a0, DISPLAY_TX_DATA
    j     poll_keyboard
```

- **Assembler.** Software that translates assembly language to machine language.
- **Machine language.** Sequence of bits that represent the most basic operations a machine may perform.

Computational abstraction

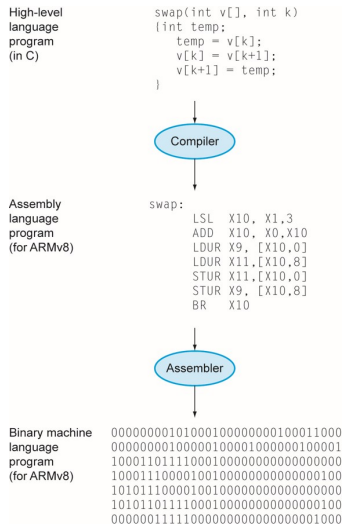
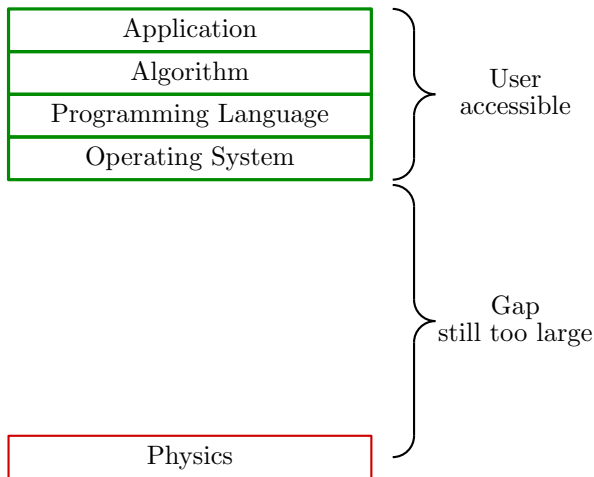


Figure 2: High-level language to machine language.

- The higher the level of the language, the more flexibility it provides.
- High-level languages allow programs to be independent of the computer on which they are developed and deployed.

Computational abstraction



Computational abstraction

- Let's now move to the physical layer.
- Do these two things represent the same functionality?

```
1 module NAND2(input logic A,  
2             input logic B,  
3             output logic X);  
4     assign X = ~(A & B);  
5 endmodule
```

Listing 1: SystemVerilog NAND2 module.




Figure 3: NAND2 layout.

Computational abstraction

- What about these representations?
- Truth table.

A	B	X
0	0	1
0	1	1
1	0	1
1	1	0

- Schematic symbol.




Figure 4: NAND2 gate.

Computational abstraction

- What about these representations?




Figure 5: NAND2 transistor level.






Figure 6: NAND2 layout.

Computational abstraction



- Even with the **user** and **technology** layers, there is a gap to bridge.
- **Q: How do we fill this gap?**
- **A: Computer Architecture.**

Computational abstraction

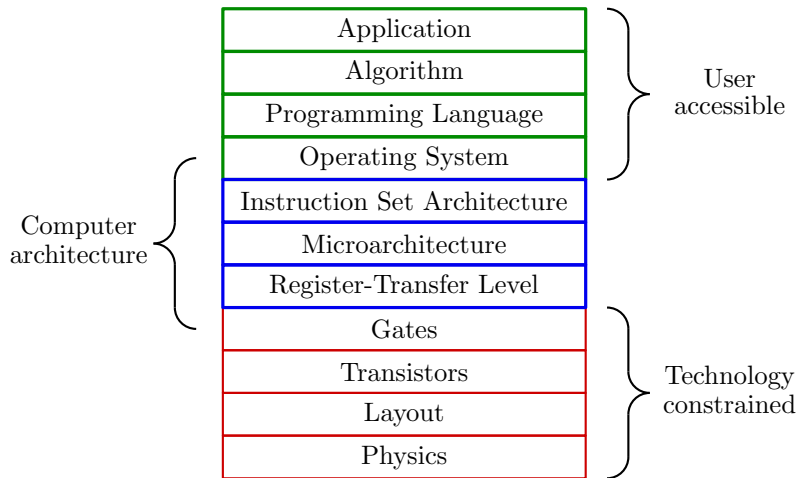


Figure 7: Simplified abstraction levels in a computing system.

- Computer architecture is the bridge between the user and the physics.
- Computer architecture may be seen as the interaction of the basic HW building blocks of a computer, as well as the semantics and rules required for such interaction.
- The three main components of computer architecture are
 - Instruction Set Architecture (ISA).
 - Microarchitecture (μA).
 - Register-Transfer Level (RTL).

Summary

- Abstraction allows us to concentrate on what really matters.
- Engineers deal with an abstraction level according to their specialization. For example:
 - SW engineers do not directly interact with HW. However, they must consider some hardware aspects such as memory allocation and usage.
 - Logic designers may deal with RTL and netlist.
 - Physical design engineers may deal with netlists and layouts.
 - Layout designers work in routing all the different semiconductor layers in an IC.
 - Analogue designers may work at the transistor and schematic levels.
- The higher the abstraction level, the more it may enclose.
- Moreover, the concept of abstraction allowed to see where this course is focused on.

Instruction Set Architecture

Instruction Set Architecture

Definition

Instruction Set Architecture (ISA)

- It is an **abstract** concept which defines the portion of a computer that is visible to both the programmer and the compiler.
- It is part of the link between an application (something a human does such as video recording, playing music, editing a spreadsheet, etc) and the physical layer of the computer, *i.e.* HW.

Instruction Set Architecture

ISA vs μA

- ISA **theoretically** describes how a computer executes its programs.
- It describes:
 - The fundamental operations, which are simply referred to as **instructions**, that the computer can execute.
 - How these instructions are executed.
 - The semantics and rules required for the interaction of the different building blocks of a computer.

Instruction Set Architecture

ISA vs μA

- Overall, ISA provides valuable information to the programmer.
 - Is a computer stack-, accumulator- or register-based?
 - Does the computer have memory? Does it have registers?
 - How many steps, *i.e.* clock cycles, does it take to execute instructions?
 - Where are operands fetched from?
 - Where is the result stored?
 - How big are data types?

Instruction Set Architecture

ISA vs μA

μA

- μA is more closely related to the **physical** implementation of a design, *i.e.*, μA determines how the ISA is implemented in HW.
 - For example, it describes which building are necessary in order to model a μP and how these building blocks are connected with each other.

Instruction Set Architecture

ISA vs μA

- The same ISA may be physically implemented in a variety of μA s.
 - For example, one ISA could be implemented by different HW approaches and vendors such as Intel, ARM or AMD, and all three could have different performances.
- A naive adder example:
 - ISA specifies data width as 64-bits.
 - μA defines the adder as ripple-carry, carry-lookahead, carry-save, carry-select, etc.

Instruction Set Architecture

ISA vs μA

- The goal of a processor designer is to evaluate the different trade-offs between ISA and μA in order to find a Pareto optimal system.
 - Power consumption.
 - Latency - How long does it take to complete a task.
 - Throughput - How many tasks can be completed in a given time.
 - Chip area.

Instruction Set Architecture

Same ISA, different μA - 45 nm technology

- x86 ISA.
- Quad Core.
- 2.6 GHz.
- 125 W.




Figure 8: AMD Phenom II X4

- x86 ISA.
- Dual Core.
- 1.6 GHz.
- 4 W.




Figure 9: Intel Atom 230

Instruction Set Architecture

Different ISA, different μ A- 45 nm technology

- x86 ISA.
- Quad Core.
- 2.6 GHz.
- 125 W.
- Power ISA.
- Octa Core.
- 4.25 GHz.
- 200 W.




Figure 10: AMD Phenom II X4




Figure 11: IBM Power7

ISA characteristics

ISA characteristics

- Type and size of instructions.
- Type and size of operands.
- Instruction encoding.
- Addressing modes.
- Registers

Registers

ISA characteristics

Registers

A register is a small memory element that stores data for quick access. This is list of registers commonly found in μ Ps and μ P-based systems. Note that generally, systems do not implement every single register in this list.

- **Program Counter (PC)**. Also called **Instruction Pointer (IP)**, points to the memory address of the next instruction to be executed.
- **Register File (RF)**. Set of registers used to store data.
- **Stack Pointer (SP)**. Points to the next location in the stack. Used in PUSH and POP operations.

Basic registers of a computer

- **Instruction Register (IR)**. Holds the instruction currently being executed.
- **Memory Address Register (MAR)**. Also called **Address Register (AR)**, points to the memory address to/from which data is stored/fetched to/from.
- **Instruction Memory (IM)**. Memory that stores the instructions that the μ P will execute.

Basic registers of a computer

- **Accumulator (ACC)**. Holds the result of arithmetic and logic operations.
- **Data Memory (DM)**. Also called **Data Register (DR)**, holds operand(s) to be used in arithmetic and logic operations.
- **General Purpose Register (GPR)**. Registers to temporary store data or addresses.
- **Status Register (SR)**. Also called **Flag Register (FR)**, holds the special conditions of the result of arithmetic and logic operations, as well as branch and jump status. For example, indicates if a comparison resulted in an equality, if the result of an operation is zero, negative, overflow, *etc.*

Basic registers of a computer

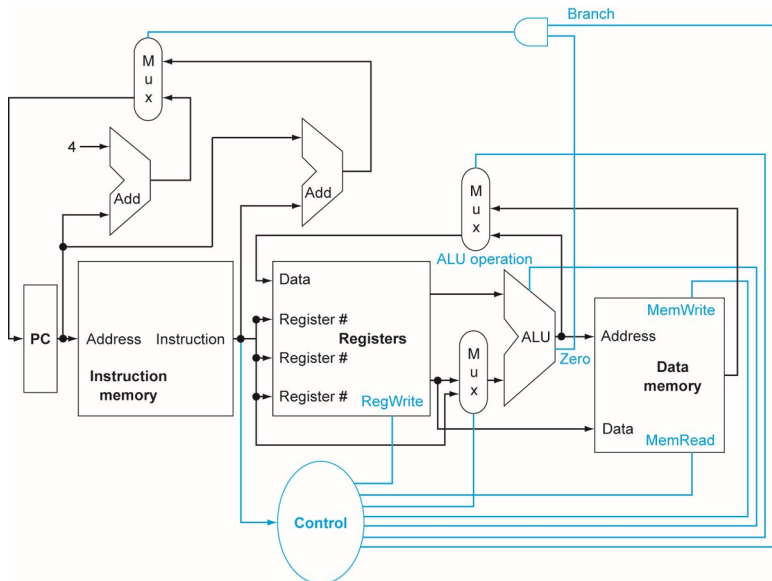


Figure 12: Basic structure of an ARM μ P-based system.

Instructions

ISA characteristics

Instructions

- Instructions are the basic operations that a μP can understand and perform.
- Programmers must use basic operations in order to complete and implement complex and high-level algorithms.
- The complete set of instructions that a μP may perform is called **instruction set**, **which is not the same as Instruction Set Architecture!**

ISA characteristics

Instructions

Instruction type	Example ¹
Arithmetic & logical	ADD, SUB, AND, OR
Data transfer	LOAD, SW, MOV, PUSH, POP
Conditional branch	BNE, BEQ
Unconditional jumps	JMP, JAL, CALL, RET
System	RD_INT, PRNT_CHR, SYSCALL
Floating Point	FADD, FMULT
String	MOVSB, STR_MV, STR_CMP
Signal processing ²	ADD_ARRAY, MULT_ARRAY, FFT

¹These examples are not specific to a particular ISA.

²Typically found in Single Instruction Multiple Data (SIMD) ISAs.

ISA characteristics

Instructions

Table 1: Intel's x86 top ten instructions based on five SPECint92 programs³.

Rank	Type	Distribution
1	load	22%
2	conditional branch	20%
3	compare	16%
4	store	12%
5	add	8%
6	and	6%
7	sub	5%
8	move register-register	4%
9	call	1%
10	return	1%
Total		96%

³J. L. Hennessy and D. A. Patterson, *Computer architecture: A quantitative approach*, 6th ed., p A-4, Morgan Kaufmann, 2019.

Operands and operations

ISA characteristics

Operands and operations

- Where do operands come from?
- Where are results stored?
- What is the size of the operands?
- How many steps does an instruction take?

ISA characteristics

Operands and operations

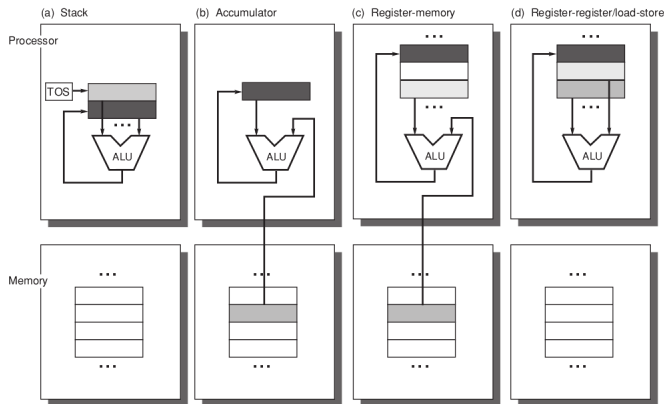



Figure 13: Operand locations for different ISAs [Figure A.1] ⁴.

⁴J. L. Hennessy and D. A. Patterson, *Computer architecture: A quantitative approach*, 6th ed., p A-4, Morgan Kaufmann, 2019.

ISA characteristics

Operands and operations



Stack-based ISA

$$C = A + B$$

Push A

Push B


Add

Pop C

ISA characteristics

Operands and operations

(b) Accumulator



Accumulator-based ISA

$$C = A + B$$

Load A


Add B

Store C

ISA characteristics

Operands and operations

(c) Register-memory



Memory-Register-based ISA


$$C = A + B$$

Load	R1	A
Add	R3	R1 B
Store	R3	C

ISA characteristics

Operands and operations

(d) Register-register/load-store



Register-Register-based ISA

$$C = A + B$$

Load	R1	A	
Load	R2	B	
Add	R3	R1	R2
Store	R3	C	

Addressing Modes

ISA characteristics

Addressing modes

- How can we read/write data from/into memory?
- What types of memory exist?

ISA characteristics

Addressing modes

Table 2: Examples of addressing modes

Mode	Example	Meaning
Immediate	ADD R4 , 3	$R4 \leftarrow R4 + 3$
Register	ADD R4 , R3	$R4 \leftarrow R4 + R3$
Absolute (Direct)	ADD R2 , (100)	$R2 \leftarrow R2 + \text{Mem}[100]$
Register indirect	ADD R4 , (R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$
Indexed	ADD R3 , (R1 + R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$
Displacement	ADD R4 , 100(R1)	$R4 \leftarrow R4 + \text{Mem}[100+R1]$
Memory indirect	ADD R1 , @(R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$

ISA characteristics

Addressing modes

Table 3: Examples of addressing modes

Mode	Example	Meaning
Autoincrement	ADD R1 , (R2)+	$R1 \leftarrow R1 + \text{Mem}[R2]$ $R2 \leftarrow R2 + d$
Autodecrement	ADD R1 , -(R2)	$R2 \leftarrow R2 - d$ $R1 \leftarrow R1 + \text{Mem}[R2]$
Scaled	ADD R1 , 100(R2) [R3]	$R1 \leftarrow R1 + \text{Mem}[100 \cdot R2 + R3 \cdot d]$

Note that Autoincrement and Autodecrement modes, the order of + and - signs influence the order of the operations. For example, -(R2) indicates decrementing R2 before accessing to memory.


ISA characteristics

Addressing modes

1. Immediate addressing




2. Register addressing



Registers

Register


3. Base addressing



Memory



+



4. PC-relative addressing



Memory



+




Figure 14: ARMv6 basic addressing modes.

ISA characteristics

Addressing modes: Example

- Let's assume that registers R_i , $i \in [1, 4]$, and selected memory locations of a computer store the following values.

Reg	Val	Mem	Val
R1	23	7	23
R2	11	11	13
R3	7	13	31
R4	19	23	17
		34	37
		100	13
		123	29
		132	41

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Immediate	ADD R4 , 3	$R4 \leftarrow R4 + 3$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Register	ADD R4 , R3	$R4 \leftarrow R4 + R3$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Absolute (Direct)	ADD R2 , (100)	$R2 \leftarrow R2 + \text{Mem}[100]$

Reg	Val	Mem	Val	R2 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Register indirect	ADD R4 , (R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Indexed	ADD R3 , (R1 + R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$

Reg	Val	Mem	Val	R3 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Displacement	ADD R4 , 100(R1)	$R4 \leftarrow R4 + \text{Mem}[100+R1]$

Reg	Val	Mem	Val	R4 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Memory indirect	ADD R1 , @(R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$

Reg	Val	Mem	Val	R1 = ?
R1	23	7	23	
R2	11	11	13	
R3	7	13	31	
R4	19	23	17	
		34	37	
		100	13	
		123	29	
		132	41	

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Autoincrement	ADD R1 , (R2)+	$R1 \leftarrow R1 + \text{Mem}[R2]$ $R2 \leftarrow R2 + d$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 3$
 $R1 = ?$

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Autodecrement	ADD R1 , -(R2)	$R2 \leftarrow R2 - d$ $R1 \leftarrow R1 + \text{Mem}[R2]$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 4$
 $R1 = ?$

ISA characteristics

Addressing modes: Example

Mode	Example	Meaning
Scaled	ADD R1 , 100(R2) [R3]	$R1 \leftarrow R1 + \text{Mem}[100+R2 + R3 *d]$

Reg	Val
R1	23
R2	11
R3	7
R4	19

Mem	Val
7	23
11	13
13	31
23	17
34	37
100	13
123	29
132	41

$d = 3$
 $R1 = ?$

Instruction encoding

ISA characteristics

Instructions encoding

- In stored-program computers, instructions and data are stored in memory.
- So, how does a processor know
 - how to differentiate between operations and operands?
 - which instruction to perform?
 - which Reg or Mem locations are the operands located?
 - which Reg or Mem locations should the result be stored to?
 - how to differentiate between

$R1 \leftarrow R1 + R1$

$R1 \leftarrow R1 + \text{Mem}[R1]$

$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R1]]$

$R1 \leftarrow \text{Mem}[R1] + \text{Mem}[R1]$

- Instruction encoding is a **convention used to differentiate the various operations in a μP , as well as operations from operands.**

ISA characteristics

Instructions encoding

- Instructions are encoded using binary representation.
- Suppose we want to design a processor that can implement the following instructions.

$R1 \leftarrow R1 + R1$

$R1 \leftarrow R1 + \text{Mem}[R1]$

$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R1]]$

$R1 \leftarrow \text{Mem}[R1] + \text{Mem}[R1]$

- We could assign a binary code to each of these 4 operations.

Instruction	Binary code
$R1 \leftarrow R1 + R1$	00
$R1 \leftarrow R1 + \text{Mem}[R1]$	01
$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R1]]$	10
$R1 \leftarrow \text{Mem}[R1] + \text{Mem}[R1]$	11

- Let's try to expand this implementation.

ISA characteristics

Instructions encoding: A naive example

- Let R_d be the destination register and R_{si} the source register, where $d \in [0, 3]$ and $i \in [0, 3]$.
- We could assign a binary code for each combination of R_d and R_{si} in the instruction $R_d \leftarrow R_{s1} + R_{s2}$.

Instruction	Binary code	Hex code
$R_0 \leftarrow R_0 + R_0$	00 0000	00h
$R_0 \leftarrow R_0 + R_1$	00 0001	01h
\vdots	\vdots	\vdots
$R_1 \leftarrow R_2 + R_3$	01 1011	1Bh
\vdots	\vdots	\vdots
$R_2 \leftarrow R_3 + R_0$	10 1100	2Ch
\vdots	\vdots	\vdots
$R_3 \leftarrow R_3 + R_3$	11 1111	3Fh

ISA characteristics

Instructions encoding: A naive example

- What about the μA of this encoding?




Figure 15: A naive μA for adding two registers.

ISA characteristics

Instructions encoding

- Is our previous scheme feasible?
- What's wrong with it?
- Could it be generalised?
- What about other addressing modes?
- How could we include other operations such as subtractions or jumps?
- Are all instructions represented using the same number of bits?

ISA characteristics

Instructions encoding

- We can continue to expand this scheme in order to add other operations, *e.g.*, subtraction and logical operations.
- Moreover, we can continue to include bits that represent different addressing modes.
- The ultimate goal of this, is to design an encoding feasible for all operations and addressing modes in our ISA.

ISA characteristics

Instructions encoding

- We could have a bit for selecting addition/subtraction in our previous design.

Table 4: Naive encoding for adding and subtracting two numbers.

Instruction	Binary code	Hex code
$R0 \leftarrow R0 + R0$	000 0000	00h
\vdots	\vdots	\vdots
$R1 \leftarrow R2 + R3$	001 1011	1Bh
\vdots	\vdots	\vdots
$R3 \leftarrow R3 + R3$	011 1111	3Fh
$R0 \leftarrow R0 - R0$	100 0000	40h
\vdots	\vdots	\vdots
$R1 \leftarrow R2 - R3$	101 1011	5Bh
\vdots	\vdots	\vdots
$R3 \leftarrow R3 - R3$	111 1111	7Fh

ISA characteristics


Instructions encoding

- As seen in the previous example, the source, destination and operation type may be represented with a **single** binary code.
- This concept may be further expanded for other operations and addressing modes.
- For this purpose, we may use longer binary codes, which may be divided into several fields.
- For example, we may use a field called **opcode** in order to represent the type of operation to be performed.
- Another field may represent the **source**, *i.e.*, the memory location where data to operate with is.
- Another field may represent the **destination**, *i.e.*, the memory location where the result of the operation should be stored to.


ISA characteristics

Instructions encoding

Variable encoding



Fixed encoding



Hybrid encoding





Figure 16: Generalised instruction encoding.

ISA characteristics

Instructions encoding

- **Variable encoding.**

- Supports any number of operands, with each operand having a specific addressing mode.
- The number of encoded bits varies between instructions.
- **Compact machine programs.**
- **Harder to decode.**

- **Fixed encoding.**

- Every instruction has the same number of operands, with addressing modes specified in the opcode.
- The number of encoded bits is always the same regardless of the instruction or addressing modes of the operands.
- **Easier to decode.**
- **Wasted bits in some instructions.**

- **Hybrid encoding.**

- Instructions use two different encoding lengths (16- and 32-bits, for example).

ISA characteristics

Instructions encoding

- ISAs may be classified into two main categories according to the complexity of their instruction encoding.
- **Reduced Instruction Set Computer (RISC).**
- **Complex Instruction Set Computer (CISC).**

CISC

- ISAs that perform complex operations and the instruction formats are not uniform.
- Large number of instructions available.
- Microcode approach.
 - A single instruction may be divided into several smaller instructions.
 - For example, a single instruction may perform a load from memory, an arithmetic operation and a store to memory.
- Reduced size of the compiled code due to **variable-length** encoding.
 - Shortest encodings represent the most commonly used instructions.

RISC

- ISAs that have a small number of simple, **fixed-length** instructions.
- Single-cycle instructions.
- Load-store approach.
 - Only load and store instructions are used for transferring data between registers and memory.

ISA characteristics

CISC vs RISC

```
1:  mul16:
2:      pushl   %ebp                ; 01010101
3:      movl    %esp, %ebp          ; 1000100111100101
4:      movl    8(%ebp), %ecx        ; 100010000100110100001000
5:      pushl   %ebx                ; 01010011
6:      movl    12(%ebp), %edx       ; 100010110101010100001100
7:      xorl    %ebx, %ebx          ; 0011000111011011
8:      movl    $15, %eax           ; 1011100000001111
9:      .p2align 2,,3               ; 000000000000000000000000
                                   ; 100011010111011000000000
10:  .L6:
11:      testb   $1, %dl             ; 111101101100001000000001
12:      je      .L5                 ; 0111010000000010
13:      addl    %ecx, %ebx          ; 0000000111001011
14:  .L5:
15:      sall    %ecx                ; 1101000111100001
16:      shrl    %edx                ; 1101000111101010
17:      decl    %eax                ; 01001000
18:      jns     .L6                 ; 0111100111110010
19:      movl    %ebx, %eax          ; 1000100111011000
20:      popl    %ebx                ; 01011011
21:      leave   ; 11001001
22:      ret      ; 11000011
```

Figure 17: CISC code example.

ISA characteristics

CISC vs RISC

```
1:  mul16:
2:      move    $6, $0                # 00000000000000000011000000100001
3:      li      $3, 15                # 00100100000000011000000000000111
4:  $L6:
5:      andi    $2, $5, 0x1           # 00110000101000100000000000000001
6:      addiu   $3, $3, -1            # 00100100011000111111111111111111
7:      beq     $2, $0, $L5           # 00010000010000000000000000000010
8:      srl     $5, $5, 1              # 00000000000001010010100001000010
9:      addu    $6, $6, $4            # 00000000110001000011000000100001
10: $L5:
11:      bgez    $3, $L6               # 00000100011000011111111111111010
12:      sll     $4, $4, 1              # 00000000000001000010000001000000
13:      j       $31                    # 00000011111000000000000000001000
14:      move    $2, $6                # 00000000110000000001000000100001
```

Figure 18: RISC code example.

QUIZ

- Which of the following are affected by the instruction encoding?
 - A) The execution time of each instruction.
 - B) The μA of the processor.
 - C) Global warming.
 - D) The size of the compiled program.
 - E) All of the above.
 - F) None of the above.

QUIZ

- Which of the following are affected by the instruction encoding?
 - A) **The execution time of each instruction.**⁵
 - B) **The μ A of the processor.**
 - C) Global warming.
 - D) **The size of the compiled program.**
 - E) All of the above.
 - F) None of the above.

⁵Think about CISC and RISC differences.

Summary

- ISA is the link between applications and HW.
- μA refers to the physical implementation of the ISA.
- The same ISA can be implemented in different μA s.
- ISA encloses
 - Type and size of instructions and operands.
 - Addressing modes.
 - Instruction encoding.
- There are RISC and CISC ISAs.
- There are several trade-offs associated between ISAs and μA s, and our goal is to find a Pareto-optimal design.

Further Reading

- Read about the difference between Von-Neumann and Harvard architectures.